

Optional Lab: Gradient Descent for Logistic Regression

Goals

In this lab, you will:

- update gradient descent for logistic regression.
- explore gradient descent on a familiar data set

```
In [1]: import copy, math
import numpy as np
%matplotlib widget
import matplotlib.pyplot as plt
from lab_utils_common import dlc, plot_data, plt_tumor_data, sigmoid, compute_cost_logistic
from plt_quad_logistic import plt_quad_logistic, plt_prob
plt.style.use('./deeplearning.mplstyle')
```

Data set

Let's start with the same two feature data set used in the decision boundary lab.


```
In [2]: X_train = np.array([[0.5, 1.5], [1, 1], [1.5, 0.5], [3, 0.5], [2, 2], [1, 2.5]])
y_train = np.array([0, 0, 0, 1, 1, 1])
```

As before, we'll use a helper function to plot this data. The data points with label $y = 1$ are shown as red crosses, while the data points with label $y = 0$ are shown as blue circles.

```
In [3]: fig, ax = plt.subplots(1, 1, figsize=(4, 4))
plot_data(X_train, y_train, ax)

ax.axis([0, 4, 0, 3.5])
ax.set_ylabel('$x_1$', fontsize=12)
ax.set_xlabel('$x_0$', fontsize=12)
plt.show()
```

Figure



Logistic Gradient Descent

Recall the gradient descent algorithm utilizes the gradient calculation:

$$\begin{aligned} &\text{repeat until convergence: } \{ \\ &\quad w_j = w_j - \alpha \frac{\partial J(\mathbf{w}, b)}{\partial w_j} \quad \text{for } j := 0..n-1 \\ &\quad b = b - \alpha \frac{\partial J(\mathbf{w}, b)}{\partial b} \\ &\} \end{aligned}$$

Where each iteration performs simultaneous updates on w_j for all j , where

Gradient descent for logistic regression

$$\begin{aligned} &\text{repeat } \{ \\ &\quad w_j = w_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m (f_{\mathbf{w}, b}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)} \right] \\ &\quad b = b - \alpha \left[\frac{1}{m} \sum_{i=1}^m (f_{\mathbf{w}, b}(\mathbf{x}^{(i)}) - y^{(i)}) \right] \\ &\} \text{ simultaneous updates} \\ &\text{Linear regression } f_{\mathbf{w}, b}(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} + b \\ &\text{Logistic regression } f_{\mathbf{w}, b}(\mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{w} \cdot \mathbf{x} + b)}} \end{aligned} \quad (1)$$

Same concepts:

- Monitor gradient descent (learning curve)
- Vectorized implementation
- Feature scaling

$$\frac{\partial J(\mathbf{w}, b)}{\partial w_j} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{\mathbf{w},b}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)} \quad (2)$$

$$\frac{\partial J(\mathbf{w}, b)}{\partial b} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{\mathbf{w},b}(\mathbf{x}^{(i)}) - y^{(i)}) \quad (3)$$

- m is the number of training examples in the data set
- $f_{\mathbf{w},b}(\mathbf{x}^{(i)})$ is the model's prediction, while $y^{(i)}$ is the target
- For a logistic regression model

$$z = \mathbf{w} \cdot \mathbf{x} + b$$

$$f_{\mathbf{w},b}(x) = g(z)$$

where $g(z)$ is the sigmoid function:

$$g(z) = \frac{1}{1+e^{-z}}$$

Gradient Descent Implementation

The gradient descent algorithm implementation has two components:

- The loop implementing equation (1) above. This is `gradient_descent` below and is generally provided to you in optional and practice labs.
- The calculation of the current gradient, equations (2,3) above. This is `compute_gradient_logistic` below. You will be asked to implement this week's practice lab.

Calculating the Gradient, Code Description

Implements equation (2),(3) above for all w_j and b . There are many ways to implement this. Outlined below is this:

- initialize variables to accumulate `dj_dw` and `dj_db`
- for each example
 - calculate the error for that example $g(\mathbf{w} \cdot \mathbf{x}^{(i)} + b) - \mathbf{y}^{(i)}$
 - for each input value $x_j^{(i)}$ in this example,
 - multiply the error by the input $x_j^{(i)}$, and add to the corresponding element of `dj_dw`. (equation 2 above)
 - add the error to `dj_db` (equation 3 above)
- divide `dj_db` and `dj_dw` by total number of examples (m)
- note that $\mathbf{x}^{(i)}$ in numpy `X[i,:]` or `X[i]` and $x_j^{(i)}$ is `X[i,j]`

```
In [9]: def compute_gradient_logistic(X, y, w, b):
    """
    Computes the gradient for logistic regression

    Args:
    X (ndarray (m,n)): Data, m examples with n features
    y (ndarray (m,)): target values
    w (ndarray (n,)): model parameters
    b (scalar)       : model parameter
    Returns
    dj_dw (ndarray (n,)): The gradient of the cost w.r.t. the parameters w.
    dj_db (scalar)       : The gradient of the cost w.r.t. the parameter b.
    """
    m,n = X.shape

    # compute hypothesis
    f_wb = sigmoid(np.dot(X,w) + b)

    # compute bias derivative
    dj_db = np.sum(f_wb-y) / m

    # compute weight derivative
    dj_dw = np.dot(np.transpose(X),(f_wb-y)) / m

    return dj_db, dj_dw
```

Check the implementation of the gradient function using the cell below.

```
In [10]: X_tmp = np.array([[0.5, 1.5], [1,1], [1.5, 0.5], [3, 0.5], [2, 2], [1, 2.5]])
y_tmp = np.array([0, 0, 0, 1, 1, 1])
w_tmp = np.array([2.,3.])
b_tmp = 1.
dj_db_tmp, dj_dw_tmp = compute_gradient_logistic(X_tmp, y_tmp, w_tmp, b_tmp)
print(f"dj_db: {dj_db_tmp}")
print(f"dj_dw: {dj_dw_tmp.tolist()}")
```

dj_db: 0.49861806546328574

dj_dw: [0.498333393278696, 0.49883942983996693]

Expected output

```
dj_db: 0.49861806546328574
dj_dw: [0.498333393278696, 0.49883942983996693]
```

Gradient Descent Code

The code implementing equation (1) above is implemented below. Take a moment to locate and compare the functions in the routine to the equations above.

```
In [11]: def gradient_descent(X, y, w_in, b_in, alpha, num_iters):
        """
        Performs batch gradient descent

        Args:
            X (ndarray (m,n)) : Data, m examples with n features
            y (ndarray (m,)) : target values
            w_in (ndarray (n,)) : Initial values of model parameters
            b_in (scalar) : Initial values of model parameter
            alpha (float) : Learning rate
            num_iters (scalar) : number of iterations to run gradient descent

        Returns:
            w (ndarray (n,)) : Updated values of parameters
            b (scalar) : Updated value of parameter
        """
        # An array to store cost J and w's at each iteration primarily for graphing Later
        J_history = []
        w = copy.deepcopy(w_in) #avoid modifying global w within function
        b = b_in

        for i in range(num_iters):
            # Calculate the gradient and update the parameters
            dj_db, dj_dw = compute_gradient_logistic(X, y, w, b)

            # Update Parameters using w, b, alpha and gradient
            w = w - alpha * dj_dw
            b = b - alpha * dj_db

            # Save cost J at each iteration
            if i<100000: # prevent resource exhaustion
                J_history.append( compute_cost_logistic(X, y, w, b) )

            # Print cost every at intervals 10 times or as many iterations if < 10
            if i% math.ceil(num_iters / 10) == 0:
                print(f"Iteration {i:4d}: Cost {J_history[-1]} ")

        return w, b, J_history #return final w,b and J history for graphing
```

Let's run gradient descent on our data set.

```
In [12]: w_tmp = np.zeros_like(X_train[0])
        b_tmp = 0.
        alph = 0.1
        iters = 10000

        w_out, b_out, _ = gradient_descent(X_train, y_train, w_tmp, b_tmp, alph, iters)
        print(f"\nupdated parameters: w:{w_out}, b:{b_out}")
```

```
Iteration    0: Cost 0.684610468560574
Iteration 1000: Cost 0.1590977666870457
Iteration 2000: Cost 0.08460064176930078
Iteration 3000: Cost 0.05705327279402531
Iteration 4000: Cost 0.04290759421682
Iteration 5000: Cost 0.03433847729884557
Iteration 6000: Cost 0.02860379802212006
Iteration 7000: Cost 0.02450156960879306
Iteration 8000: Cost 0.02142370332569295
Iteration 9000: Cost 0.019030137124109114
```


```
updated parameters: w:[5.28 5.08], b:-14.222409982019837
```

Let's plot the results of gradient descent:

```
In [13]: fig,ax = plt.subplots(1,1,figsize=(5,4))
        # plot the probability
        plt_prob(ax, w_out, b_out)

        # Plot the original data
        ax.set_ylabel(r'$x_1$')
        ax.set_xlabel(r'$x_0$')
        ax.axis([0, 4, 0, 3.5])
        plot_data(X_train,y_train,ax)

        # Plot the decision boundary
        x0 = -b_out/w_out[0]
        x1 = -b_out/w_out[1]
        ax.plot([0,x0],[x1,0], c='d1blue', lw=1)
        plt.show()
```



In the plot above:

- the shading reflects the probability $y=1$ (result prior to decision boundary)
- the decision boundary is the line at which the probability = 0.5


Another Data set

Let's return to a one-variable data set. With just two parameters, w , b , it is possible to plot the cost function using a contour plot to get a better idea of what gradient descent is up to.

```
In [14]: x_train = np.array([0., 1, 2, 3, 4, 5])
         y_train = np.array([0, 0, 0, 1, 1, 1])
```

As before, we'll use a helper function to plot this data. The data points with label $y = 1$ are shown as red crosses, while the data points with label $y = 0$ are shown as blue circles.

```
In [15]: fig, ax = plt.subplots(1,1,figsize=(4,3))
         plt_tumor_data(x_train, y_train, ax)
         plt.show()
```

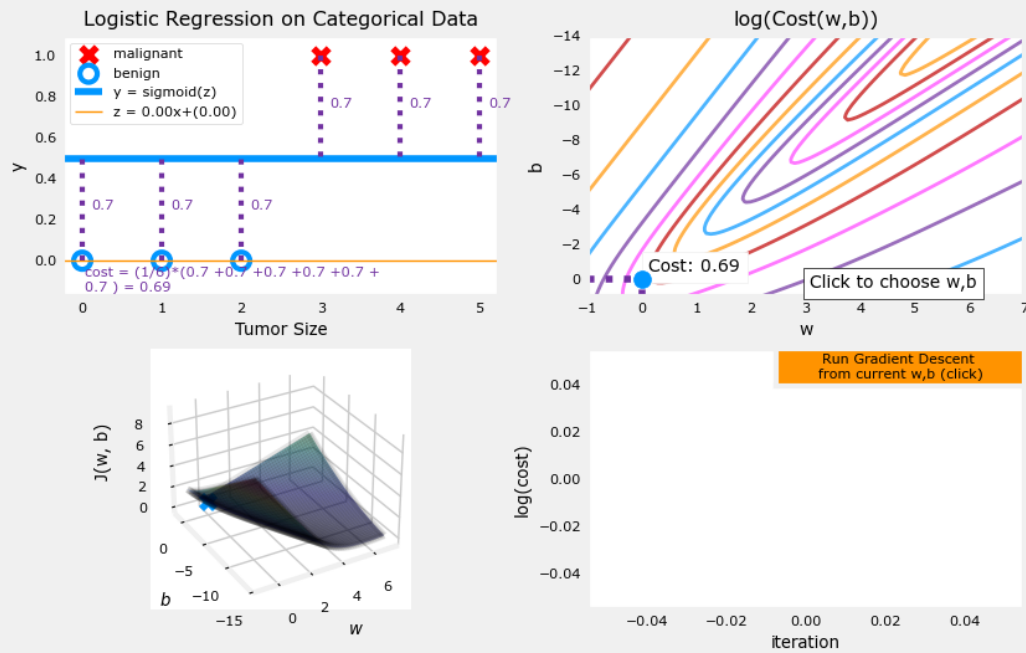


In the plot below, try:

- changing w and b by clicking within the contour plot on the upper right.
 - changes may take a second or two
 - note the changing value of cost on the upper left plot.
 - note the cost is accumulated by a loss on each example (vertical dotted lines)
- run gradient descent by clicking the orange button.
 - note the steadily decreasing cost (contour and cost plot are in $\log(\text{cost})$)
 - clicking in the contour plot will reset the model for a new run
- to reset the plot, rerun the cell

```
In [16]: w_range = np.array([-1, 7])
         b_range = np.array([1, -14])
         quad = plt_quad_logistic( x_train, y_train, w_range, b_range )
```

Figure



Congratulations!

You have:

- examined the formulas and implementation of calculating the gradient for logistic regression
- utilized those routines in
 - exploring a single variable data set
 - exploring a two-variable data set